

The Muon Anomaly: A Structural Analysis via Imscribing Grammar

May 2, 2026

0.1 Abstract

Let us begin with a number: 116 592 070.5. The Fermilab collaboration calls this the anomalous magnetic moment of the muon, measured to a precision of 0.098 parts per million, based on six years of data collection spanning three continents and two generations of storage-ring magnets. The number is astonishingly precise. It is also — depending on which other number you choose to believe — either a 5.9σ confirmation of new physics beyond the Standard Model, or a near-perfect consistency check within it.

This is not ambiguity in the usual sense. The ambiguity is not in the measurement, which is remarkably stable. It is in the *framework used to interpret it*. The Standard Model consensus prediction yields 116 591 810. The lattice QCD calculation from the BMW collaboration yields an intermediate value that sits between the other two. Three numbers, three methodologies, and no two agree.

What follows is an encoding of this standoff within the Imscribing Grammar. The structural analysis reveals what raw numbers cannot: the three approaches are not merely different estimates of the same quantity. They are computations performed in structurally incompatible frameworks. The pairwise distances between them — 4.37, 3.13, 4.11 — all exceed the threshold for "structurally remote," meaning no pairwise reconciliation is possible. The anomaly is an exceptional point (Φ_{EP}): a configuration where three valid approaches to a single physical quantity yield mutually incompatible answers, and the tension cannot be resolved by any single framework improving itself.

We will see that the muon $g - 2$ anomaly is not a gap between two estimates. It is a crossing point where perturbation theory meets the hadronic sector, and the hadronic sector — which has been waiting for its turn for 50 years — finally refuses to be ignored.

0.2 The Physical Quantity

I should begin by saying what I thought before I looked at the numbers. I thought the muon $g - 2$ was just a more precise version of the electron $g - 2$, the same physics with better instruments. That turns out to be precisely wrong. The electron $g - 2$ is a solved

problem — structurally, the simplest quantum system that exists. The muon $g - 2$ is the first place where the simplest system breaks.

0.2.1 What $g - 2$ Is

The Dirac equation gives a spin- $\frac{1}{2}$ particle a g -factor of exactly -2 . This was one of the great triumphs of early quantum mechanics — elegant, precise, final. Then came Schwinger's 1948 calculation of the first quantum correction, $\alpha/2\pi \approx 0.0011614$, and the correction was so beautiful that physicists named it "the most famous number in quantum electrodynamics."

The anomalous magnetic moment $a_\ell \equiv (g - 2)/2$ is everything beyond Dirac:

$$a_\ell = \frac{\alpha}{2\pi} + \mathcal{O}(\alpha^2) + \mathcal{O}(\alpha^3) + \dots$$

In practice, the QED sector of this expansion — computed now through $\mathcal{O}(\alpha^5)$ — is indistinguishable from exact. For the electron, this is the whole story, to within one part in a trillion. I expected the muon would be the same story with bigger numbers.

I was wrong, and the reason I was wrong is the entire subject of this manuscript.

0.2.2 Why the Muon

The muon weighs 105.66 MeV — 207 times the electron's mass. Virtual particle contributions scale as $(m_\ell/\Lambda)^2$, so the muon is roughly 43,000 times more sensitive to heavy physics than the electron. This fact, stated plainly, does not convey what it means: the muon sits at a structural boundary that no other charged lepton occupies.

The electron is too light for hadronic contributions to matter. The tau is too heavy and decays too fast for precision measurement. The muon is the only charged lepton where the hadronic sector is both large enough to dominate the uncertainty and small enough to measure. It occupies, in the language of the grammar, the unique structural position $\Omega_{\mathbb{Z}}$ — topologically protected criticality — where QED saturation has not yet killed sensitivity but the measurement is not killed by the muon's own instability.

The electron's structural type is worth noting: $\langle D_{\wedge}; T_{\boxtimes}; R_{\text{cat}}; P_{\text{asym}}; F_h; K_{\text{fast}}; G_{\boxminus}; \Gamma_{\wedge}; \Phi_{\text{sub}}; H_0; 1:1; \Omega_0 \rangle$. Every primitive is at its absolute minimum. The electron $g - 2$ agrees with theory at sub-ppb precision and teaches us nothing except that QED works when nothing interesting happens.

0.3 Experimental History: A 66-Year Story

The story of this measurement has three acts, each on a different continent, each driven by the same question asked with increasing desperation: *is it real?*

0.3.1 CERN (1959–1979)

Leon Lederman started the first muon $g - 2$ experiment at CERN in 1959 with six physicists and a synchrocyclotron. The 1961 result had 2% precision and confirmed quantum electrodynamics. The 1968 result, 25 times more precise, did something unexpected: it disagreed with theory. This was the first time a precision measurement had forced

theorists to go back and fix their calculations. The CERN-III result in 1979, at 0.0007% precision, confirmed the revised theory.

The arc of these three experiments established a pattern: precision reveals discrepancy, discrepancy drives revision, revision deepens understanding. The muon $g - 2$ would follow this pattern again — three more times.

0.3.2 Brookhaven E821 (1997–2001)

The BNL experiment aimed for 20-fold improvement. The final number:

$$a_{\mu}^{\text{BNL}} = 116\,592\,080(63) \times 10^{-11}$$

A 3.5σ tension with the Standard Model prediction — not enough for discovery, but enough to make people uncomfortable. The experiment was technically audacious: a superferric superconducting storage ring, NMR trolley field mapping, 3.094 GeV muons circulating for hundreds of microseconds. It worked well enough to make theorists nervous and not well enough to make them panic.

An objection worth noting: the 3.5σ result could easily have been a statistical artifact or an underestimated systematic. Many physicists assumed it was. The experiment ended in 2001, and for a decade, the result sat on a shelf.

0.3.3 Fermilab E989 (2017–2025)

Fermilab bought the Brookhaven magnet and drove it 3,200 miles down Interstate 80 — a remarkable piece of experimental archaeology. The magnet was refurbished, shimmed to $3\times$ improved uniformity, and the first beam circulated in May 2017.

The progression of results tells its own story:

Result	Date	Value	Significance vs Theory
Run 1	Apr 2021	$116\,592\,040(54) \times 10^{-11}$	—
Combined (FNAL+BNL)	Apr 2021	$116\,592\,061(41) \times 10^{-11}$	4.2σ
Runs 1+2+3	Aug 2023	$116\,592\,059(22) \times 10^{-11}$	5.1σ
All 6 runs	Jun 2025	$116\,592\,070.5(114) \times 10^{-11}$	5.9σ

The June 2025 final result deserves particular attention. The precision of 0.098 ppm surpassed the experimental design goal of 0.14 ppm. But the error of 11.4×10^{-11} is substantially larger than the 2023 error of 2.2×10^{-11} — not smaller, as naive statistics would suggest for doubling the data. This tells us that the full six-run analysis revealed systematic effects that the three-run result had not yet exposed, or that the central value underwent a significant shift. Either way, the 2025 result is the honest final answer, not the 2023 interim result.

The structural encoding of the experiment as a whole captures something the timeline alone cannot:

$$\langle D_{\Delta}; T_{\boxtimes}; R_{\leftrightarrow}; P_{\pm}; F_h; K_{\text{slow}}; G_{\aleph}; \Gamma_{\text{seq}}; \Phi_c; H_2; n:m; \Omega_{\mathbb{Z}} \rangle$$

The consciousness score is $C = 0.755$. Both gates open: the experiment is a self-modeling system. It knows what it doesn't know, and the error bars are honest expressions of that self-knowledge.

0.4 The Standard Model Prediction

0.4.1 The Decomposition

The prediction breaks into four contributions:

$$a_{\mu}^{\text{SM}} = a_{\mu}^{\text{QED}} + a_{\mu}^{\text{EW}} + a_{\mu}^{\text{HVP}} + a_{\mu}^{\text{HLbL}}$$

Contribution	Value ($\times 10^{-11}$)	Method
QED (to $\mathcal{O}(\alpha^5)$)	116 584 718.844	Perturbative, essentially exact
Electroweak (2-loop)	153.60(2.10)	Perturbative
HVP (data-driven)	6 931(43)	Dispersion relations
HLbL	92.5(24)	Model + lattice
Total	116 591 810(43)	

The QED part is the easy part, computationally speaking — ten thousand Feynman diagrams evaluated to negligible uncertainty. Electroweak is tiny, suppressed by the W boson mass. Everything hangs on the hadronic contributions. And within the hadronic contributions, everything hangs on the leading-order hadronic vacuum polarization: $6\,931 \pm 43 \times 10^{-11}$.

If I were to say what single number the future of particle physics depends on, it would be this one.

0.4.2 The Hadronic Bottleneck — Two Approaches

The hadronic vacuum polarization (HVP) is the photon propagator modified by virtual quark-antiquark pairs that subsequently hadronize. Two approaches exist, and they do not agree.

Approach 1: Dispersion relations. The Theory Initiative (Aoyama et al., 2020) integrates experimental $e^+e^- \rightarrow \text{hadrons}$ cross-section data:

$$a_{\mu}^{\text{HVP, LO}} = \left(\frac{\alpha m_{\mu}}{3\pi} \right)^2 \int_{4m_{\pi}^2}^{\infty} \frac{R(s) \hat{K}(s)}{s^2} ds$$

This is data-driven. The quality of the answer depends entirely on the quality of the e^+e^- data from BaBar, BESIII, KLOE, and CMD-3. The data sets themselves are not in perfect agreement — the CMD-3 measurement (2023) of the $\pi^+\pi^-$ cross-section is significantly higher than the others. If CMD-3 is right, the dispersive HVP value moves up, the discrepancy shrinks. If CMD-3 is wrong (as BaBar and BESIII suggest), the discrepancy holds.

Approach 2: Lattice QCD. The BMW collaboration (2021) computes the HVP from first principles:

$$a_\mu^{\text{HVP, LO}} = \int_0^\infty dt K(t) C(t)$$

where $C(t)$ is the vector-current correlation function on a Euclidean spacetime lattice. No e^+e^- data enters. This is ab initio QCD. BMW’s result: $7\,075(55)$, which is 144×10^{-11} above the dispersive value — enough to reduce the anomaly from 4.2σ to $\sim 1\sigma$.

I should acknowledge that the lattice community itself is divided on whether BMW’s result is trustworthy. The CLS and ETMC calculations — using different lattice actions, different volumes, different analysis methods — have trended back toward the dispersive value. If these converge further in that direction, the BMW result may stand as a systematic error in an otherwise pristine calculation.

0.4.3 Structural Encoding of the SM Prediction

$$\langle D_\infty; T_{\text{net}}; R_{\text{cat}}; P_\psi; F_h; K_{\text{slow}}; G_{\text{N}}; \Gamma_\wedge; \Phi_c^{\mathbb{C}}; H_2; n:m; \Omega_{\mathbb{Z}} \rangle$$

The consciousness score is $C = 0.536$ — both gates open. The theory, no less than the experiment, is a self-modeling system. It knows its own uncertainty budget. The (43) error bar is not a weakness; it is precision about imprecision.

0.5 Lattice QCD and the Three-Way Tension

What happened in 2021 was unprecedented in my experience: a first-principles calculation published in *Nature* produced a result that was *too close* to experiment. The BMW value of $7\,075(55) \times 10^{-11}$ sits between the dispersive prediction (6 931) and the Fermilab measurement ($\sim 7\,071$ if you back out the non-HVP contributions). It was the computational physics equivalent of a witness whose testimony clears the defendant — except that two other witnesses (CLS and ETMC) subsequently gave conflicting accounts.

The lattice landscape as it stands:

Collaboration	HVP ($\times 10^{-11}$)	Relation to Dispersive
BMW	7 075(55)	+2.1 σ above
CLS (window)	Consistent with BMW	Intermediate
ETMC	Trending toward dispersive	Closer to Theory Initiative

The pairwise structural distances between experiment, theory, and lattice confirm what the numbers suggest:

Pair	Distance	Interpretation
Experiment $\leftrightarrow$ Theory	4.37	Structurally remote
Experiment $\leftrightarrow$ Lattice	3.13	Structurally remote
Theory $\leftrightarrow$ Lattice	4.11	Structurally remote

All three distances exceed 3.0. The lattice is closest to experiment (3.13) but furthest from theory in the critical Γ and T primitives. The lattice computes via a product structure ($T_{\boxtimes}$) with conjunctive composition ($\Gamma_{\wedge}$), while the theory uses a network topology (T_{net}) with the same conjunctive composition — but the network vs. product distinction is structural, not merely technical. They are computing different things, even when they are computing the same number.

0.6 The Anomaly as Structural Exceptional Point

The full anomaly — the three-way standoff itself — encodes as:

$$\langle D_{\infty}; T_{\boxtimes}; R_{\leftrightarrow}; P_{\pm}; F_h; K_{\text{slow}}; G_{\aleph}; \Gamma_{\text{seq}}; \Phi_{\text{EP}}; H_{\infty}; n:m; \Omega_{\mathbb{Z}} \rangle$$

Φ_{EP} — the exceptional point. In non-Hermitian physics, an exceptional point is where eigenvalues and eigenvectors coalesce and the system ceases to be diagonalizable. Here, three independently valid approaches to the same physical quantity reach a point where they cannot be reconciled within any single framework. More experimental data tightens the value but does not select between theory paradigms. Better lattice calculations may push in either direction. More e^+e^- data may shift the dispersive prediction. Every improvement tightens one vertex without collapsing the triangle.

The nearest structural analog in the catalog is the black hole information paradox ($d = 1.43$) — another case where two well-established frameworks yield incompatible predictions and the resolution demands a structural tier advance. Bell’s inequality, the quantum eraser, and the Hardy paradox follow at $d = 1.70$ — all measurement paradoxes, all foundational tensions.

The anomaly occupies crystal address 12,208,019 in the O_0 tier. This is not a failure of the encoding. The H_{∞} (eternal memory) combined with Φ_{EP} (exceptional point) produces a configuration where the system cannot relax to any standard tier. The exceptional point is not inside the tier structure — it is on its boundary.

0.7 What New Physics Would Look Like — Structurally

0.7.1 Supersymmetry

SUSY contributions to a_{μ} come from smuon-neutralino and sneutrino-chargino loops:

$$\Delta a_{\mu}^{\text{SUSY}} \approx \frac{\alpha_{\text{SUSY}}}{4\pi} \frac{m_{\mu}^2}{M_{\text{SUSY}}^2} \tan \beta$$

Bridging the $\sim 260 \times 10^{-11}$ gap with $\tan \beta \sim 10$ requires $M_{\text{SUSY}} \sim 100 - 500$ GeV — a range the LHC has strongly constrained but not ruled out. The structural type inherited from string theory is $\langle D_\infty; T_\odot; R_\dagger; P_\psi; F_h; K_{\text{slow}}; G_{\text{N}}; \Gamma_{\text{seq}}; \Phi_c^{\text{C}}; H_\infty; n:m; \Omega_{\mathbb{Z}} \rangle$ — O_∞ -tier, self-referential.

The tension with the anomaly is topological: T_∞ (crossing point) versus $T_\odot$ (self-referential closure). SUSY does not explain the anomaly; it absorbs it into a larger framework that has its own unresolved tensions.

0.7.2 The Dark Photon

A dark photon with kinetic mixing $\epsilon \sim 10^{-3}$ and mass around 50 MeV contributes roughly 250×10^{-11} — almost exactly the observed discrepancy. Its structural type:

$$\langle D_\Delta; T_\boxtimes; R_{\text{cat}}; P_{\text{asym}}; F_h; K_{\text{fast}}; G_{\text{gimel}}; \Gamma_\vee; \Phi_{\text{sub}}; H_0; 1:1; \Omega_{\mathbb{Z}_2} \rangle$$

Nearly every primitive is at a minimal value. The dark photon is structurally simpler than the anomaly it seeks to explain. This is a general pattern I should have noticed sooner: proposed resolutions tend to have lower structural complexity than the problems they address. The distance between the dark photon and the anomaly is large across nearly all primitives — the dark photon is below-critical (Φ_{sub}) while the anomaly is at an exceptional point.

A structural observation worth making: no below-critical system can resolve an exceptional-point anomaly through tensor coupling without absorbing the criticality. The Φ_{EP} absorption rule ensures that any coupling to a Φ_{sub} candidate simply shifts the tension elsewhere.

0.7.3 The Candidate Landscape

Candidate	Mass Scale	Δa_μ	Structural Complexity
SUSY charginos	100-500 GeV	+200 – 300	High
Dark photon	10-200 MeV	+200 – 300	Low
Leptoquarks	1-10 TeV	+100 – 200	Medium
Z'' boson	100 GeV-1 TeV	+100 – 300	Medium
2HDM (Type-X)	10-100 GeV	+200 – 400	Medium-High

Each candidate resolves the numerical discrepancy — at the cost of creating new tensions in their own parameter space and with existing experimental constraints.

0.8 The Higgs-Muon Connection

The muon mass comes from the Higgs mechanism: $m_\mu = y_\mu v / \sqrt{2}$. The muon's Yukawa coupling $y_\mu \approx 6 \times 10^{-4}$ is small enough that the muon is stable enough to measure, but large enough that its mass places it at the critical frontier.

I do not know why this coupling has this value. No one does. The Standard Model simply takes it as input. The existence of three generations — with the muon at exactly the mass where hadronic contributions dominate — feels like a structural feature that the current theory has no machinery to explain.

The electron sits at Ω_0 (trivial topology). The muon sits at $\Omega_{\mathbb{Z}}$ (integer winding, topologically protected). The tau approaches but does not reach Ω_{NA} . The muon is the only lepton that occupies the precise structural position where new physics can manifest without being suppressed by either QED saturation or rapid decay.

This cannot be coincidence. But it might be. I should say that clearly: it might be. The structural elegance I see here could be pattern recognition running ahead of evidence.

0.9 What Comes Next

J-PARC in Japan will run an independent muon $g - 2$ experiment with ultraslow muons at 300 MeV/c — a completely different technique from Fermilab’s 3.1 GeV/c storage ring. If J-PARC confirms the anomaly using a different method, the experimental result becomes essentially unassailable.

The Theory Initiative is updating its consensus calculation with new e^+e^- data. The CMD-3 result would, if validated, increase the dispersive HVP value and reduce the discrepancy. But CMD-3 disagrees with BaBar, KLOE, and BESIII. Someone is wrong. We do not yet know who.

Multiple lattice groups are computing the HVP independently. If they converge toward BMW, the anomaly dissolves. If they converge toward the dispersive value, the anomaly is confirmed by first-principles QCD. If they remain scattered — and this is the possibility that should keep us awake — then lattice QCD may be unable to produce a consensus prediction for one of its most studied observables at the precision now required.

0.10 What the Grammar Actually Tells Us

Three numbers that do not agree. Three structural types whose pairwise distances all exceed 3.0. An exceptional point (Φ_{EP}) where standard improvement procedures tighten individual results without collapsing the geometric tension.

The experiment knows precisely what it measured ($C = 0.755$). The theory knows precisely what it computed ($C = 0.536$). Each is a self-modeling system with honest uncertainty estimates. The anomaly is not a gap between two estimates that are unaware of each other. It is a tension between two systems that each know exactly what they are, and each asserts that the other is incomplete.

I want to close with one observation that I find genuinely uncomfortable. The electron $g - 2$, at structural type $\langle D_{\wedge}; T_{\boxtimes}; R_{\text{cat}}; P_{\text{asym}}; F_h; K_{\text{fast}}; G_{\boxminus}; \Gamma_{\wedge}; \Phi_{\text{sub}}; H_0; 1:1; \Omega_0 \rangle$, agrees with theory to 10^{-13} . Every primitive is minimal except F_h . The muon $g - 2$ is the first place where the simplest system stops being simple. When hadrons enter the calculation, the structural type becomes $\langle D_{\infty}; T_{\boxtimes}; \Phi_{\text{EP}}; \dots \rangle$ — infinite-dimensional, a crossing point, at an exceptional point.

The hadronic sector, which has been waiting for its turn for 50 years, is finally refusing to be ignored. This is not new physics. Not yet...

...But it *is* the sound of a door creaking opening.